



# Gradient **DESCENT**

- In Stats, ML and other Data Science field, We optimize a lot of stuff.
- In Linear Regression, we optimize intercept and slope. Eq. :  $\text{y} = \text{mx} + \text{c}$
- In Logistic one, we optimize a squiggle.


When we use **Logistic Regression**, we optimize a squiggle.



And when we use **t-SNE**, we optimize clusters.




- Gradient Descent can optimize all this stuff and much more.
- Suppose, you have 3 points in a graph for which you have to find optimized intercept and slope.
- First, You assume intercept so that you have something to improve upon. (Also, slope as least square estimate)




- Find residual between observed and predicted and using that → Sum of Squared Residuals.  

$$\sum_{i=1}^n (\text{observed}_i - \text{predicted}_i)^2$$
- Then plot a graph of Intercept(X) - Sum of Squared Residuals(Y)

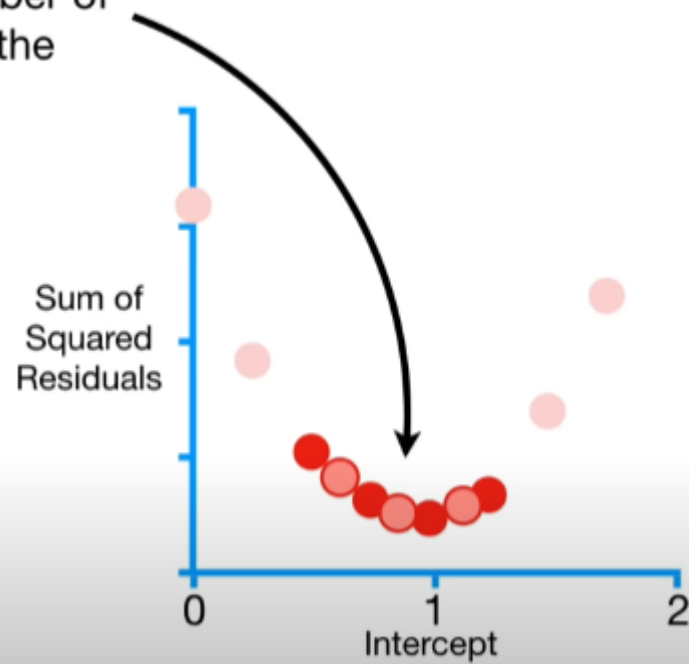
Of the points that we calculated for the graph, this one has the lowest Sum of Squared Residuals...




**Gradient Descent** only does a few calculations far from the optimal solution...




...and increases the number of calculations closer to the optimal value.



...and baby steps when it is close.



...then calculated the **New Intercept**,  
the difference between the **Old Intercept** and the **Step Size**.



**Step 1:** Take the derivative of the **Loss Function** for each parameter in it. In fancy Machine Learning Lingo, take the **Gradient** of the **Loss Function**.

**Step 2:** Pick random values for the parameters.

**Step 3:** Plug the parameter values into the derivatives (ahem, the **Gradient**).

**Step 4:** Calculate the Step Sizes: **Step Size** = **Slope** × **Learning Rate**

**Step 5:** Calculate the New Parameters:

$$\text{New Parameter} = \text{Old Parameter} - \text{Step Size}$$

Now go back to **Step 3** and repeat until **Step Size** is very small, or you reach the **Maximum Number of Steps**.

**Step 3:** Plug the parameter values into the derivatives (ahem, the **Gradient**).

**Step 4:** Calculate the Step Sizes: **Step Size** = **Slope** × **Learning Rate**

**Step 5:** Calculate the New Parameters:

$$\text{New Parameter} = \text{Old Parameter} - \text{Step Size}$$

- For large datasets, this approach could be in-efficient.

So there is a thing called  
**Stochastic Gradient Descent**  
that uses a randomly selected  
subset of the data at every step  
rather than the full dataset.

This reduces the time spent  
calculating the derivatives of the  
**Loss Function.**